

An Efficient Representation for Lazy Constructors using 64-bit Pointers

Georgios Fourtounis (gfour@softlab.ntua.gr)
Nikolaos Papaspyrou (nickie@softlab.ntua.gr)



National Technical University of Athens
School of Electrical and Computer Engineering
Software Engineering Laboratory

The 3rd ACM SIGPLAN Workshop on
Functional High-Performance Computing (FHPC'14)
Göteborg, Sweden, September 4, 2014

Work supported by the project Handling Uncertainty in Data Intensive Applications, co-financed by the European Union (European Social Fund — ESF) and Greek national funds, through the Operational Program "Education and Lifelong Learning", under the program THALES.



In case you fall asleep...

What the **paper** is about

An efficient implementation technique based on tagged pointers on AMD64, used in a compiler for a subset of Haskell

In case you fall asleep...

What the **paper** is about

An efficient implementation technique based on tagged pointers on AMD64, used in a compiler for a subset of Haskell

What this **talk** will be about

In case you fall asleep...

What the **paper** is about

An efficient implementation technique based on tagged pointers on AMD64, used in a compiler for a subset of Haskell

What this **talk** will be about



In case you fall asleep...

What the **paper** is about

An efficient implementation technique based on tagged pointers on AMD64, used in a compiler for a subset of Haskell

What this **talk** will be about

- ① An outline of the techniques used in the **gic** compiler that enable this optimization on AMD64

In case you fall asleep...

What the **paper** is about

An efficient implementation technique based on tagged pointers on AMD64, used in a compiler for a subset of Haskell

What this **talk** will be about

- ① An outline of the techniques used in the **gic** compiler that enable this optimization on AMD64
 - The intensional transformation

In case you fall asleep...

What the **paper** is about

An efficient implementation technique based on tagged pointers on AMD64, used in a compiler for a subset of Haskell

What this **talk** will be about

- ① An outline of the techniques used in the **gic** compiler that enable this optimization on AMD64
 - The intensional transformation
 - Lazy activation records

In case you fall asleep...

What the **paper** is about

An efficient implementation technique based on tagged pointers on AMD64, used in a compiler for a subset of Haskell

What this **talk** will be about

- ① An outline of the techniques used in the **gic** compiler that enable this optimization on AMD64
 - The intensional transformation
 - Lazy activation records
- ② `flip filter thePaper $
liftA2 (&&) isInteresting canBePresentedHere`

The intensional transformation

Alternative technique for implementing **non-strict functional languages** by transformation to dataflow programs

- [Yaghi, 1984] The intensional implementation technique for functional languages.
- [Arvind & Nikhil, 1990] The “coloring” technique for implementing functions on the MIT Dataflow Machine.
- [Rondogiannis & Wadge, 1997, 1999] A formalization of the intensional transformation and its extension for a class of higher-order programs.

Dataflow programming languages

Dataflow programming:

- A program is a directed graph of **data** flowing through a network of **processing units**
- Quite popular in the 1980s due to its implicitly parallel nature

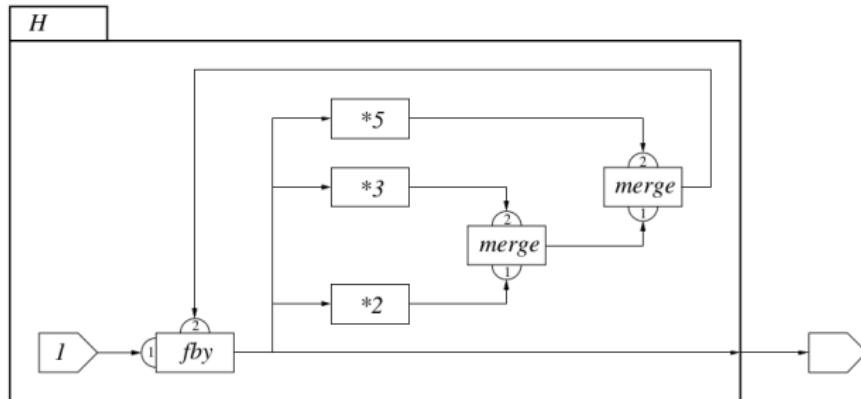


Figure from Joey Paquet's PhD thesis, "Intensional Scientific Programming" (1999)

Dataflow programming languages

Dataflow programming:

- A program is a directed graph of **data** flowing through a network of **processing units**
- Quite popular in the 1980s due to its implicitly parallel nature

Dataflow languages:

- Mostly **functional** in nature, encouraging **stream processing**
- **Examples:** Val, Id, Lucid, GLU, SISAL, etc.

Dataflow programming languages

Dataflow programming:

- A program is a directed graph of **data** flowing through a network of **processing units**
- Quite popular in the 1980s due to its implicitly parallel nature

Dataflow languages:

- Mostly **functional** in nature, encouraging **stream processing**
- **Examples:** Val, Id, Lucid, GLU, SISAL, etc.

Dataflow machines:

- **Specialized** parallel architectures for executing dataflow programs, e.g. the MIT Tagged-Token Machine
- Execution is determined by the **availability** of input arguments to operations

The status of dataflow

In the 1990s:

- Interest started to decline
- Dataflow architectures  vs. mainstream uniprocessors 

The status of dataflow

In the 1990s:

- Interest started to decline
- Dataflow architectures  vs. mainstream uniprocessors 

Today:

- Renewed interest
- Uniprocessors no longer follow Moore's law
- Commodity parallel hardware on the rise
- A new generation of dataflow-esque languages/programming models: Dryad, Cluster, Hyrax, Map-Reduce, etc.
- Efficient implementation in mainstream **multi-core** architectures and reconfigurable hardware (FPGAs)

Intensional transformation by example

The input is a first-order functional program. The output is a program with parameterless definitions (intensional program).

Example

```
result  =  f 3 + f 5
f x      =  g (x*x)
g y      =  y+2
```

Intensional transformation by example

The input is a first-order functional program. The output is a program with parameterless definitions (intensional program).

Example

```
result  =  f 3 + f 5  
f x      =  g (x*x)  
g y      =  y+2
```

Step 1: for all functions f

- Replace the i -th call of f by $\text{call}_i(f)$
- Remove formal parameters from function definitions

Intensional transformation by example

The input is a first-order functional program. The output is a program with parameterless definitions (intensional program).

Example

result = f 3 + f 5	result = call ₀ (f)+call ₁ (f)
f x = g (x*x)	f = call ₀ (g)
g y = y+2	g = y+2

Step 1: for all functions f

- Replace the i -th call of f by $\text{call}_i(f)$
- Remove formal parameters from function definitions

Intensional transformation by example

The input is a first-order functional program. The output is a program with parameterless definitions (intensional program).

Example

result = f 3 + f 5	result = call ₀ (f)+call ₁ (f)
f x = g (x*x)	f = call ₀ (g)
g y = y+2	g = y+2

Step 2: for all functions f , for all formal parameters x

- Find actual parameters corresponding to x in all calls of f
- Introduce a new definition for x with an **actuals** clause, listing the actual parameters in the order of the calls

Intensional transformation by example

The input is a first-order functional program. The output is a program with parameterless definitions (intensional program).

Example

result = f 3 + f 5	result = call ₀ (f)+call ₁ (f)
f x = g (x*x)	f = call ₀ (g)
g y = y+2	g = y+2

x = actuals(3, 5)
y = actuals(x*x)

Step 2: for all functions f , for all formal parameters x

- Find actual parameters corresponding to x in all calls of f
- Introduce a new definition for x with an **actuals** clause, listing the actual parameters in the order of the calls

Semantics of the target language

Evaluation of expressions: $\text{EVAL}(e, w)$

- **Intensional**: with respect to a **context** w
- Evaluation contexts are **lists** of natural numbers
- The **initial** context is the empty list

Semantics of the target language

Evaluation of expressions: $EVAL(e, w)$

- **Intensional**: with respect to a **context** w
- Evaluation contexts are **lists** of natural numbers
- The **initial** context is the empty list

Context switching: call and actuals

$$\begin{aligned} EVAL(\text{call}_i(e), w) &= EVAL(e, i : w) \\ EVAL(\text{actuals}(e_0, \dots, e_{n-1}), i : w) &= EVAL(e_i, w) \end{aligned}$$

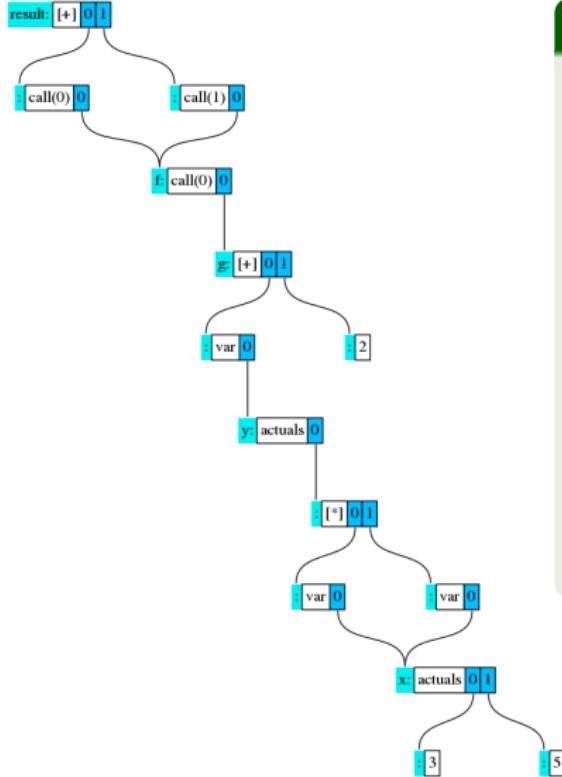
Example

Evaluation of the target program:

```
EVAL(result, [])
= EVAL(call0(f)+ call1(f), [])
= EVAL(call0(f), []) + EVAL(call1(f), [])
= EVAL(f, [0]) + EVAL(f, [1])
= EVAL(call0(g), [0]) + EVAL(call0(g), [1])
= EVAL(g, [0, 0]) + EVAL(g, [0, 1])
= EVAL(y, [0, 0]) + EVAL(2, [0, 0]) + EVAL(y, [0, 1]) + EVAL(2, [0, 1])
= EVAL(actuals(x*x), [0, 0]) + 2 + EVAL(actuals(x*x), [0, 1]) + 2
= EVAL(x*x, [0]) + 2 + EVAL(x*x, [1]) + 2
= EVAL(x, [0]) * EVAL(x, [0]) + 2 + EVAL(x, [1]) * EVAL(x, [1]) + 2
= EVAL(actuals(3, 5), [0]) * EVAL(actuals(3, 5), [0]) + 2 +
EVAL(actuals(3, 5), [1]) * EVAL(actuals(3, 5), [1]) + 2
= EVAL(3, []) * EVAL(3, []) + 2 + EVAL(5, []) * EVAL(5, []) + 2
= 3 * 3 + 2 + 5 * 5 + 2
= 38
```

```
result = call0(f)+call1(f)
f      = call0(g)
g      = y+2
x      = actuals(3, 5)
y      = actuals(x*x)
```

Dataflow graph



Example

$$\text{result} = \text{f}(3) + \text{f}(5)$$

$$\text{f } x = \text{g}(x*x)$$

$$\text{g } y = y+2$$

$$\text{result} = \text{call}_0(\text{f}) + \text{call}_1(\text{f})$$

$$\text{f} = \text{call}_0(\text{g})$$

$$\text{g} = y+2$$

$$x = \text{actuals}(3, 5)$$

$$y = \text{actuals}(x*x)$$

Higher-order functions: Multiple dimensions

The idea

- A program with functions of order up to n is translated to an intensional program using n -dimensional contexts
- Dimension i has its own pair of call^i and actuals^i operators
- Transformation works in steps, each time decreasing the program's order by one

Example (A second order program that can be transformed)

```
result      = double inc 3
double f x = f (f x)
inc y       = y + 1
```

Missing pieces

The original intensional transformation lacks:

- ① User-defined data structures:

```
data List = Nil | Cons Int List  
length ls =  
  case ls of  
    Nil          → 0  
    Cons x xs   → 1 + length xs
```

Missing pieces

The original intensional transformation lacks:

- ① User-defined data structures:

```
data List = Nil | Cons Int List  
length ls =  
  case ls of  
    Nil          → 0  
    Cons x xs   → 1 + length xs
```

- ② Partial application:

```
result      = double (add 1) 3  
double f x = f (f x)  
add a b     = a + b
```

Missing pieces

The original intensional transformation lacks:

- ① User-defined data structures:

```
data List = Nil | Cons Int List  
length ls =  
  case ls of  
    Nil          → 0  
    Cons x xs  → 1 + length xs
```

- ② Partial application:

```
result      = double (add 1) 3  
double f x = f (f x)  
add a b     = a + b
```

👍 Problem (2) reduced to (1) with **defunctionalization**

The complete picture: The **gic** compiler

<https://github.com/gfour/gic>

Key ideas:

- Compiles a large subset of **Haskell 98**
- **Modular defunctionalization** handles higher-order programs, supporting separate compilation
[Fourtounis, Papaspyrou & Theofilopoulos, 2014]
- **Generalized intensional transformation** handles first-order programs with user-defined data types
[Fourtounis, Papaspyrou & Rondogiannis, 2013]
- Currently supercombinator-based, using lambda-lifting
- Implementation on mainstream hardware using an adaptation of **lazy activation records**
[Charalambidis, Grivas, Papaspyrou & Rondogiannis, 2008]

Implementation using lazy activation records (LARs)

- Target code is **low-level C**:
 - can be compiled using gcc, clang, icc
 - uses GNU extension: statement expressions
- Each Haskell function becomes a C function taking a LAR
- Each Haskell constructor c becomes a C function taking a LAR l and returning the pair (c, l)
- **Push/enter model**: Push LAR, then enter function
- Memory use:
 - the C stack is used for control (function entry-return)
 - LARs are usually placed on the heap, but a cheap analysis enables their placement on the stack

The anatomy of a LAR (i)

Characteristics:

- LARs efficiently implement **intensional contexts** and are also used for storing actual parameters
- LARs are **lazy**, i.e. the actual parameters are filled in upon demand
- LAR construction is controlled by the presence of intensional operators in the target program; destruction may require **GC**

The anatomy of a LAR (i)

Characteristics:

- LARs efficiently implement **intensional contexts** and are also used for storing actual parameters
- LARs are **lazy**, i.e. the actual parameters are filled in upon demand
- LAR construction is controlled by the presence of intensional operators in the target program; destruction may require **GC**

Contents:

Each LAR contains three fields:

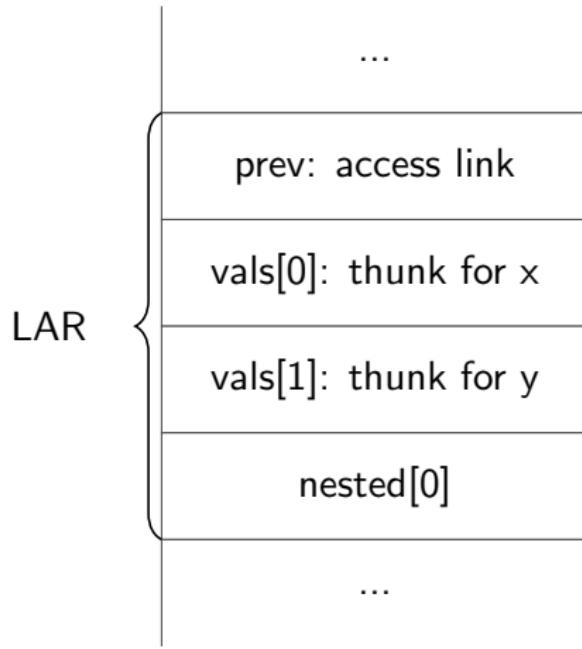
`prev` access link to parent LAR

`vals` thunks for function parameters

`nested` escaping LARs reached through pattern-matching,
used to evaluate lazy constructor fields

The anatomy of a LAR (ii)

```
f x y = case x of [] -> [1]
                  a:as -> [a + y]
```



Representation of thunks:

A thunk must have three fields:

- flag set if the thunk is evaluated

- code pointer to code that must be run in order to evaluate the thunk

- value the already computed value of an evaluated thunk

The anatomy of a LAR (iii)

Representation of thunks:

A thunk must have three fields:

flag set if the thunk is evaluated

code pointer to code that must be run in order to evaluate the thunk

value the already computed value of an evaluated thunk

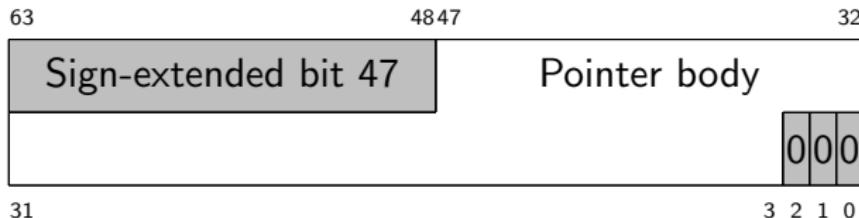
Representation of values:

A value can be:

- a primitive integer n (for enumerations, e.g. Bool and Int)
- a lazy constructor of the form (c, l)

AMD64 pointers

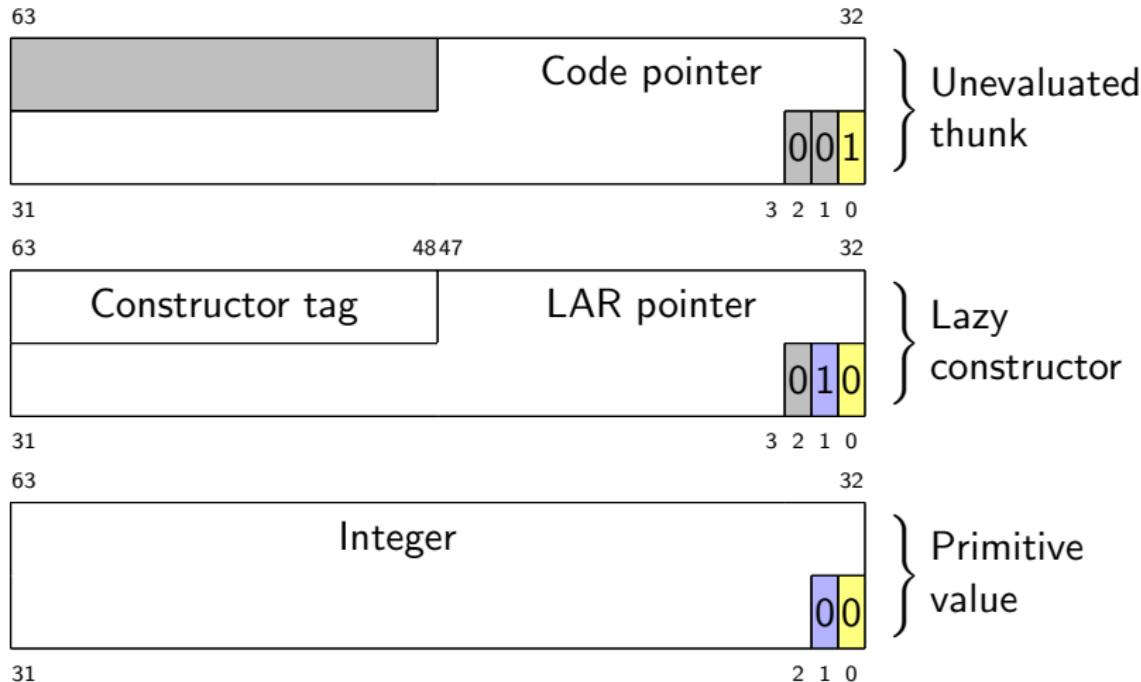
Pointers in AMD64 leave a lot of space unused:



Remarks:

- This is true for pointers to the **heap** (using a suitable memory allocator, e.g. `malloc` or our own) and to the **stack**
- Code pointers can also be aligned (`gcc -falign-functions`)
- This permits a tagged-pointer implementation for **thunks**

Compact thunk representation



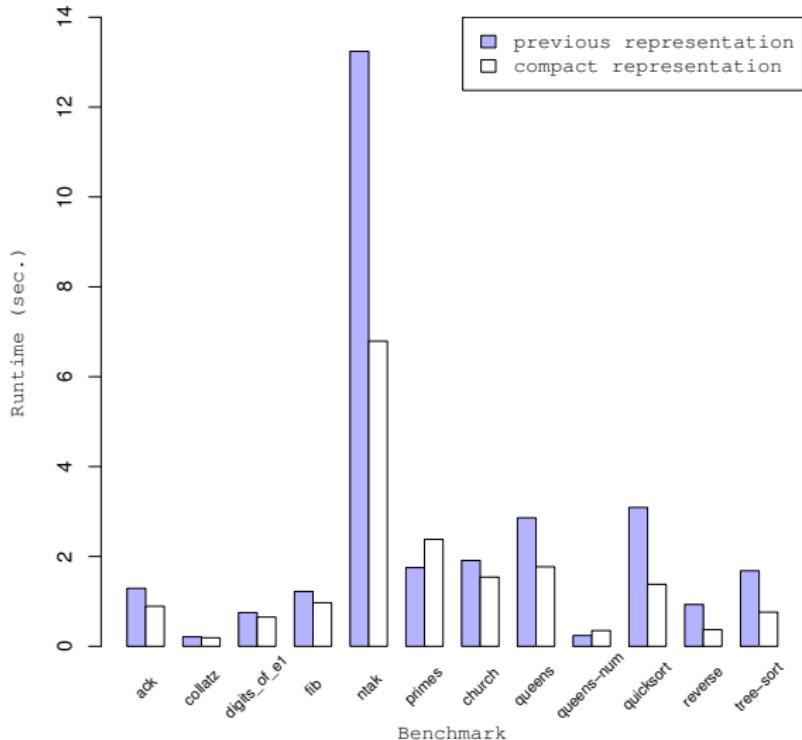
□ Thunk data ■ Thunk flag □ Value flag ■ Unused

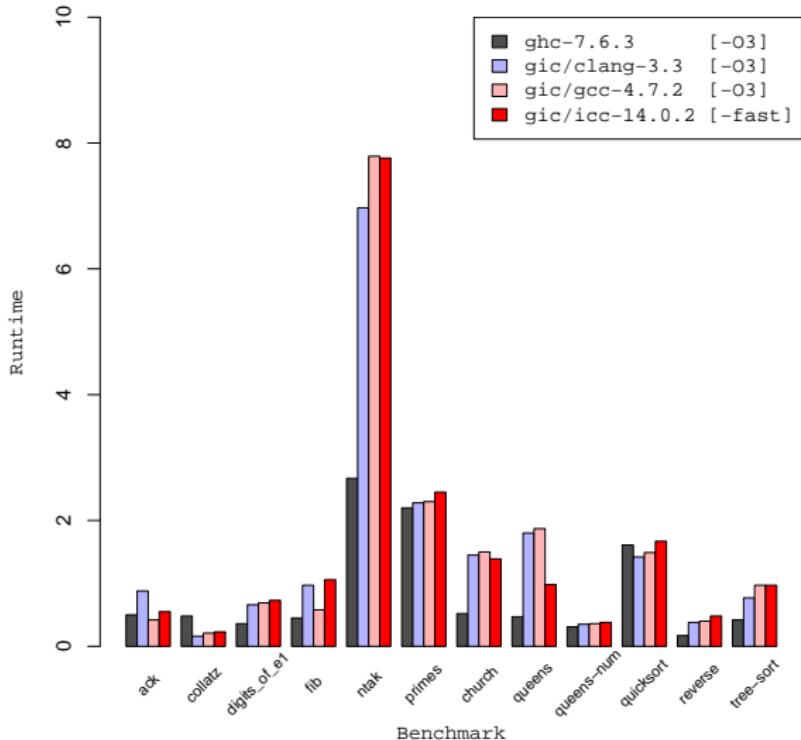
What we're comparing

- Comparison with a previous, non-compact representation
- Comparison with fully optimizing **ghc** (-O3)
 - using the default **ghc** back-end
 - using three different **gic** back-ends: gcc, clang and icc

What we're comparing

- Comparison with a previous, non-compact representation
- Comparison with fully optimizing **ghc** (-O3)
 - using the default **ghc** back-end
 - using three different **gic** back-ends: gcc, clang and icc
- We only use two basic variants of well-known optimizations (escape analysis, usage analysis), garbage collection was used
- (Micro/mini)-benchmarks, since we still don't have a front-end for full Haskell (e.g. complex patterns, type classes, full library)





So, in case you just woke up...

Summary:

- Outline of **gic**, an alternative compiler for a subset of Haskell
- Efficient implementation for **gic** with tagged pointers on AMD64

So, in case you just woke up...

Summary:

- Outline of **gic**, an alternative compiler for a subset of Haskell
- Efficient implementation for **gic** with tagged pointers on AMD64

Concluding remarks:

- Promising performance; many optimizations still missing: code transformations (e.g. fusion, inlining), tail-call optimization, strictness analysis
- Overhead on math-heavy computations, pointers pass through the ALU
- Compact memory use ($\sim 67\%$ reduction)
- The technique can also be used with ARM, SPARC processors

So, in case you just woke up...

Work in progress:

- The explicit pointer stack is expensive, working on a DWARF-based garbage collector for minimal overhead
- Tail-call optimization: also useful for defunctionalization; delicate interplay between C and non-strictness
- Add type classes using the concretization encoding of Pottier and Gauthier
- Take out the lambda-lifter and support local definitions
- Provide **gic** as a back-end for **ghc**

So, in case you just woke up...

Work in progress:

- The explicit pointer stack is expensive, working on a DWARF-based garbage collector for minimal overhead
- Tail-call optimization: also useful for defunctionalization; delicate interplay between C and non-strictness
- Add type classes using the concretization encoding of Pottier and Gauthier
- Take out the lambda-lifter and support local definitions
- Provide **gic** as a back-end for **ghc**

Future work:

- More optimizations
- Parallel runtime



Example

Evaluation of the target program:

EVAL(result, [])

```
result = call0(f)+call1(f)
f      = call0(g)
g      = y+2
x      = actuals(3, 5)
y      = actuals(x*x)
```

Example

Evaluation of the target program:

$$\begin{aligned} & \text{EVAL}(\text{result}, []) \\ = & \text{EVAL}(\text{call}_0(\text{f}) + \text{call}_1(\text{f}), []) \end{aligned}$$

```
result = call0(f)+call1(f)
f      = call0(g)
g      = y+2
x      = actuals(3, 5)
y      = actuals(x*x)
```

Example

Evaluation of the target program:

$$\begin{aligned} & \text{EVAL(result, [])} \\ = & \text{EVAL(call}_0(\text{f}) + \text{call}_1(\text{f}), []) \\ = & \text{EVAL(call}_0(\text{f}), []) + \text{EVAL(call}_1(\text{f}), []) \end{aligned}$$

result = call₀(f)+call₁(f)
f = call₀(g)
g = y+2
x = actuals(3, 5)
y = actuals(x*x)

Example

Evaluation of the target program:

$$\begin{aligned}& \text{EVAL}(\text{result}, []) \\&= \text{EVAL}(\text{call}_0(f) + \text{call}_1(f), []) \\&= \text{EVAL}(\text{call}_0(f), []) + \text{EVAL}(\text{call}_1(f), []) \\&= \text{EVAL}(f, [0]) + \text{EVAL}(f, [1])\end{aligned}$$

result = call₀(f)+call₁(f)
f = call₀(g)
g = y+2
x = actuals(3, 5)
y = actuals(x*x)

Example

Evaluation of the target program:

$$\begin{aligned}& \text{EVAL}(\text{result}, []) \\&= \text{EVAL}(\text{call}_0(\text{f}) + \text{call}_1(\text{f}), []) \\&= \text{EVAL}(\text{call}_0(\text{f}), []) + \text{EVAL}(\text{call}_1(\text{f}), []) \\&= \text{EVAL}(\text{f}, [0]) + \text{EVAL}(\text{f}, [1]) \\&= \text{EVAL}(\text{call}_0(\text{g}), [0]) + \text{EVAL}(\text{call}_0(\text{g}), [1])\end{aligned}$$

result = call₀(f)+call₁(f)
f = call₀(g)
g = y+2
x = actuals(3, 5)
y = actuals(x*x)

Example

Evaluation of the target program:

$$\begin{aligned} & \text{EVAL(result, [])} \\ = & \text{EVAL(call}_0(\text{f}) + \text{call}_1(\text{f}), []) \\ = & \text{EVAL(call}_0(\text{f}), []) + \text{EVAL(call}_1(\text{f}), []) \\ = & \text{EVAL(f, [0])} + \text{EVAL(f, [1])} \\ = & \text{EVAL(call}_0(\text{g}), [0]) + \text{EVAL(call}_0(\text{g}), [1]) \\ = & \text{EVAL(g, [0, 0])} + \text{EVAL(g, [0, 1])} \end{aligned}$$

result = call₀(f)+call₁(f)
f = call₀(g)
g = y+2
x = actuals(3, 5)
y = actuals(x*x)

Example

Evaluation of the target program:

```
EVAL(result, [ ])
= EVAL(call0(f)+ call1(f), [ ])
= EVAL(call0(f), [ ]) + EVAL(call1(f), [ ])
= EVAL(f, [0]) + EVAL(f, [1])
= EVAL(call0(g), [0]) + EVAL(call0(g), [1])
= EVAL(g, [0, 0]) + EVAL(g, [0, 1])
= EVAL(y, [0, 0]) + EVAL(2, [0, 0]) + EVAL(y, [0, 1]) + EVAL(2, [0, 1])
```

```
result = call0(f)+call1(f)
f      = call0(g)
g      = y+2
x      = actuals(3, 5)
y      = actuals(x*x)
```

Example

Evaluation of the target program:

$$\begin{aligned} & \text{EVAL(result, [])} \\ = & \text{EVAL(call}_0(\text{f}) + \text{call}_1(\text{f}), []) \\ = & \text{EVAL(call}_0(\text{f}), []) + \text{EVAL(call}_1(\text{f}), []) \\ = & \text{EVAL(f, [0])} + \text{EVAL(f, [1])} \\ = & \text{EVAL(call}_0(\text{g}), [0]) + \text{EVAL(call}_0(\text{g}), [1]) \\ = & \text{EVAL(g, [0, 0])} + \text{EVAL(g, [0, 1])} \\ = & \text{EVAL(y, [0, 0])} + \text{EVAL(2, [0, 0])} + \text{EVAL(y, [0, 1])} + \text{EVAL(2, [0, 1])} \\ = & \text{EVAL(actuals(x*x), [0, 0])} + 2 + \text{EVAL(actuals(x*x), [0, 1])} + 2 \end{aligned}$$

result = call₀(f)+call₁(f)
f = call₀(g)
g = y+2
x = actuals(3, 5)
y = actuals(x*x)

Example

Evaluation of the target program:

$$\begin{aligned} & \text{EVAL(result, [])} \\ = & \text{EVAL(call}_0(\text{f}) + \text{call}_1(\text{f}), []) \\ = & \text{EVAL(call}_0(\text{f}), []) + \text{EVAL(call}_1(\text{f}), []) \\ = & \text{EVAL(f, [0])} + \text{EVAL(f, [1])} \\ = & \text{EVAL(call}_0(\text{g}), [0]) + \text{EVAL(call}_0(\text{g}), [1]) \\ = & \text{EVAL(g, [0, 0])} + \text{EVAL(g, [0, 1])} \\ = & \text{EVAL(y, [0, 0])} + \text{EVAL(2, [0, 0])} + \text{EVAL(y, [0, 1])} + \text{EVAL(2, [0, 1])} \\ = & \text{EVAL(actuals(x*x), [0, 0])} + 2 + \text{EVAL(actuals(x*x), [0, 1])} + 2 \\ = & \text{EVAL(x*x, [0])} + 2 + \text{EVAL(x*x, [1])} + 2 \end{aligned}$$

result = call₀(f)+call₁(f)
f = call₀(g)
g = y+2
x = actuals(3, 5)
y = actuals(x*x)

Example

Evaluation of the target program:

$$\begin{aligned} & \text{EVAL(result, [])} \\ = & \text{EVAL(call}_0(\text{f}) + \text{call}_1(\text{f}), []) \\ = & \text{EVAL(call}_0(\text{f}), []) + \text{EVAL(call}_1(\text{f}), []) \\ = & \text{EVAL(f, [0])} + \text{EVAL(f, [1])} \\ = & \text{EVAL(call}_0(\text{g}), [0]) + \text{EVAL(call}_0(\text{g}), [1]) \\ = & \text{EVAL(g, [0, 0])} + \text{EVAL(g, [0, 1])} \\ = & \text{EVAL(y, [0, 0])} + \text{EVAL(2, [0, 0])} + \text{EVAL(y, [0, 1])} + \text{EVAL(2, [0, 1])} \\ = & \text{EVAL(actuals(x*x), [0, 0])} + 2 + \text{EVAL(actuals(x*x), [0, 1])} + 2 \\ = & \text{EVAL(x*x, [0])} + 2 + \text{EVAL(x*x, [1])} + 2 \\ = & \text{EVAL(x, [0])} * \text{EVAL(x, [0])} + 2 + \text{EVAL(x, [1])} * \text{EVAL(x, [1])} + 2 \end{aligned}$$

result = call₀(f)+call₁(f)
f = call₀(g)
g = y+2
x = actuals(3, 5)
y = actuals(x*x)

Example

Evaluation of the target program:

$$\begin{aligned} & \text{EVAL(result, [])} \\ = & \text{EVAL(call}_0(\text{f}) + \text{call}_1(\text{f}), []) \\ = & \text{EVAL(call}_0(\text{f}), []) + \text{EVAL(call}_1(\text{f}), []) \\ = & \text{EVAL(f, [0])} + \text{EVAL(f, [1])} \\ = & \text{EVAL(call}_0(\text{g}), [0]) + \text{EVAL(call}_0(\text{g}), [1]) \\ = & \text{EVAL(g, [0, 0])} + \text{EVAL(g, [0, 1])} \\ = & \text{EVAL(y, [0, 0])} + \text{EVAL(2, [0, 0])} + \text{EVAL(y, [0, 1])} + \text{EVAL(2, [0, 1])} \\ = & \text{EVAL(actuals(x*x), [0, 0])} + 2 + \text{EVAL(actuals(x*x), [0, 1])} + 2 \\ = & \text{EVAL(x*x, [0])} + 2 + \text{EVAL(x*x, [1])} + 2 \\ = & \text{EVAL(x, [0])} * \text{EVAL(x, [0])} + 2 + \text{EVAL(x, [1])} * \text{EVAL(x, [1])} + 2 \\ = & \text{EVAL(actuals(3, 5), [0])} * \text{EVAL(actuals(3, 5), [0])} + 2 + \\ & \text{EVAL(actuals(3, 5), [1])} * \text{EVAL(actuals(3, 5), [1])} + 2 \end{aligned}$$

result = call₀(f)+call₁(f)
f = call₀(g)
g = y+2
x = actuals(3, 5)
y = actuals(x*x)

Example

Evaluation of the target program:

```
EVAL(result, [ ])
= EVAL(call0(f)+ call1(f), [ ])
= EVAL(call0(f), [ ]) + EVAL(call1(f), [ ])
= EVAL(f, [0]) + EVAL(f, [1])
= EVAL(call0(g), [0]) + EVAL(call0(g), [1])
= EVAL(g, [0, 0]) + EVAL(g, [0, 1])
= EVAL(y, [0, 0]) + EVAL(2, [0, 0]) + EVAL(y, [0, 1]) + EVAL(2, [0, 1])
= EVAL(actuals(x*x), [0, 0]) + 2 + EVAL(actuals(x*x), [0, 1]) + 2
= EVAL(x*x, [0]) + 2 + EVAL(x*x, [1]) + 2
= EVAL(x, [0]) * EVAL(x, [0]) + 2 + EVAL(x, [1]) * EVAL(x, [1]) + 2
= EVAL(actuals(3, 5), [0]) * EVAL(actuals(3, 5), [0]) + 2 +
EVAL(actuals(3, 5), [1]) * EVAL(actuals(3, 5), [1]) + 2
= EVAL(3, [ ]) * EVAL(3, [ ]) + 2 + EVAL(5, [ ]) * EVAL(5, [ ]) + 2
```

result =	call ₀ (f)+call ₁ (f)
f =	call ₀ (g)
g =	y+2
x =	actuals(3, 5)
y =	actuals(x*x)

Example

Evaluation of the target program:

$$\begin{aligned} & \text{EVAL(result, [])} \\ = & \text{EVAL(call}_0(\text{f}) + \text{call}_1(\text{f}), []) \\ = & \text{EVAL(call}_0(\text{f}), []) + \text{EVAL(call}_1(\text{f}), []) \\ = & \text{EVAL(f, [0])} + \text{EVAL(f, [1])} \\ = & \text{EVAL(call}_0(\text{g}), [0]) + \text{EVAL(call}_0(\text{g}), [1]) \\ = & \text{EVAL(g, [0, 0])} + \text{EVAL(g, [0, 1])} \\ = & \text{EVAL(y, [0, 0])} + \text{EVAL(2, [0, 0])} + \text{EVAL(y, [0, 1])} + \text{EVAL(2, [0, 1])} \\ = & \text{EVAL(actuals(x*x), [0, 0])} + 2 + \text{EVAL(actuals(x*x), [0, 1])} + 2 \\ = & \text{EVAL(x*x, [0])} + 2 + \text{EVAL(x*x, [1])} + 2 \\ = & \text{EVAL(x, [0])} * \text{EVAL(x, [0])} + 2 + \text{EVAL(x, [1])} * \text{EVAL(x, [1])} + 2 \\ = & \text{EVAL(actuals(3, 5), [0])} * \text{EVAL(actuals(3, 5), [0])} + 2 + \\ & \text{EVAL(actuals(3, 5), [1])} * \text{EVAL(actuals(3, 5), [1])} + 2 \\ = & \text{EVAL(3, [])} * \text{EVAL(3, [])} + 2 + \text{EVAL(5, [])} * \text{EVAL(5, [])} + 2 \\ = & 3 * 3 + 2 + 5 * 5 + 2 \end{aligned}$$

result = call₀(f)+call₁(f)
f = call₀(g)
g = y+2
x = actuals(3, 5)
y = actuals(x*x)

Example

Evaluation of the target program:

$$\begin{aligned} & \text{EVAL(result, [])} \\ = & \text{EVAL(call}_0(\text{f}) + \text{call}_1(\text{f}), []) \\ = & \text{EVAL(call}_0(\text{f}), []) + \text{EVAL(call}_1(\text{f}), []) \\ = & \text{EVAL(f, [0])} + \text{EVAL(f, [1])} \\ = & \text{EVAL(call}_0(\text{g}), [0]) + \text{EVAL(call}_0(\text{g}), [1]) \\ = & \text{EVAL(g, [0, 0])} + \text{EVAL(g, [0, 1])} \\ = & \text{EVAL(y, [0, 0])} + \text{EVAL(2, [0, 0])} + \text{EVAL(y, [0, 1])} + \text{EVAL(2, [0, 1])} \\ = & \text{EVAL(actuals(x*x), [0, 0])} + 2 + \text{EVAL(actuals(x*x), [0, 1])} + 2 \\ = & \text{EVAL(x*x, [0])} + 2 + \text{EVAL(x*x, [1])} + 2 \\ = & \text{EVAL(x, [0])} * \text{EVAL(x, [0])} + 2 + \text{EVAL(x, [1])} * \text{EVAL(x, [1])} + 2 \\ = & \text{EVAL(actuals(3, 5), [0])} * \text{EVAL(actuals(3, 5), [0])} + 2 + \\ & \text{EVAL(actuals(3, 5), [1])} * \text{EVAL(actuals(3, 5), [1])} + 2 \\ = & \text{EVAL(3, [])} * \text{EVAL(3, [])} + 2 + \text{EVAL(5, [])} * \text{EVAL(5, [])} + 2 \\ = & 3 * 3 + 2 + 5 * 5 + 2 \\ = & 38 \end{aligned}$$

result = call₀(f)+call₁(f)
f = call₀(g)
g = y+2
x = actuals(3, 5)
y = actuals(x*x)

Garbage collection

- Representation compatible with a semi-space garbage collector
- LARs are the only memory objects managed, used for both function calls and constructors
- The **prev** field is a pointer:
 - Each LAR keep its layout description in the high bits of **prev**
 - The field is reused for forwarded addresses
- The **root set** contains LARs passed as arguments to functions still running
 - We currently use an explicit pointer stack (a level of indirection) to keep track of the roots without traversing the C stack
 - This results in at least 2x overhead over the no GC case (but see work in progress at the end of the talk)

Garbage collection interface

